

Q1:

When a process calls `fork()`, the operating system needs to create a child process with the same address space as the parent. If the OS immediately copied every single page from the parent's address space into new physical pages for the child, the `fork` operation would become very slow, especially for large programs.

Copy-On-Write (COW) avoids this upfront cost. Instead of duplicating all physical pages right away, the OS lets the parent and child share the exact same physical pages after the `fork`. These shared pages are marked read-only and flagged as COW pages. Nothing is actually copied yet.

The important part is what happens when one of the processes actually tries to write to one of these shared pages. Since the page is read-only, this triggers a write-fault. At that moment, the OS allocates a new physical page, copies the data from the old shared page, and then gives the writing process its own private, writable copy. The other process continues using the original page.

This technique improves `fork()` latency because:

- The OS doesn't copy the whole address space at `fork` time.
- Most child processes call `exec()` right after `fork`, so they never modify the old pages—meaning zero memory copying happens.
- Pages are only copied if they are actually written to, so the system only performs work when necessary.

As a result, COW makes `fork()` extremely fast, even if the parent process uses a lot of memory.

Sources used for Q1:

1. OS Textbook, Chapter 5
<https://pages.cs.wisc.edu/~remzi/OSTEP/cpu-api.pdf>
2. Linux Kernel Documentation – Understanding the Linux Virtual Memory Manager
<https://www.kernel.org/doc/gorman/html/understand/understand013.html>

Q2:

fork() using Copy-On-Write

```
function fork(parent):  
  
    create a new child page table
```

```
    for each page P in the parent:
```

```
        // Child shares the same physical frame as the parent  
        child.PTE[P] = parent.PTE[P]
```

```
        // Both PTEs become read-only and marked COW
```

```
        parent.PTE[P].writable = false  
        child.PTE[P].writable = false  
        parent.PTE[P].cow = true  
        child.PTE[P].cow = true
```

```
        // Increase the reference count for that frame
```

```
        refcount[P] += 1
```

```
    end for
```

```
    return child
```

```
end function
```

Handling a write to a COW page

```
on write_page_fault(process, page P):
```

```

if P is a COW page:

    old_frame = P.frame

    // Allocate a new physical page

    new_frame = allocate_frame()

    // Copy data over

    copy old_frame → new_frame

    // Update the process's page table to point to the new frame

    P.frame = new_frame

    P.writable = true

    P.cow = false

    // Decrement refcount on the old page

    refcount[old_frame] -= 1

    if refcount[old_frame] == 0:

        free(old_frame)

    end if

end if

end handler

```

Freeing pages on process exit

```
function exit(process):
```

for each page P:

 refcount[P.frame] -= 1

 if refcount[P.frame] == 0:

 free(P.frame)

 end if

end for

destroy(process)

end function